



Search

Write



Member-only story

# Practical Kubernetes. Deploy User Management Microservice.



Konstantin Mogilevskii · [Follow](#)

Published in Dev Genius · 8 min read · 4 days ago



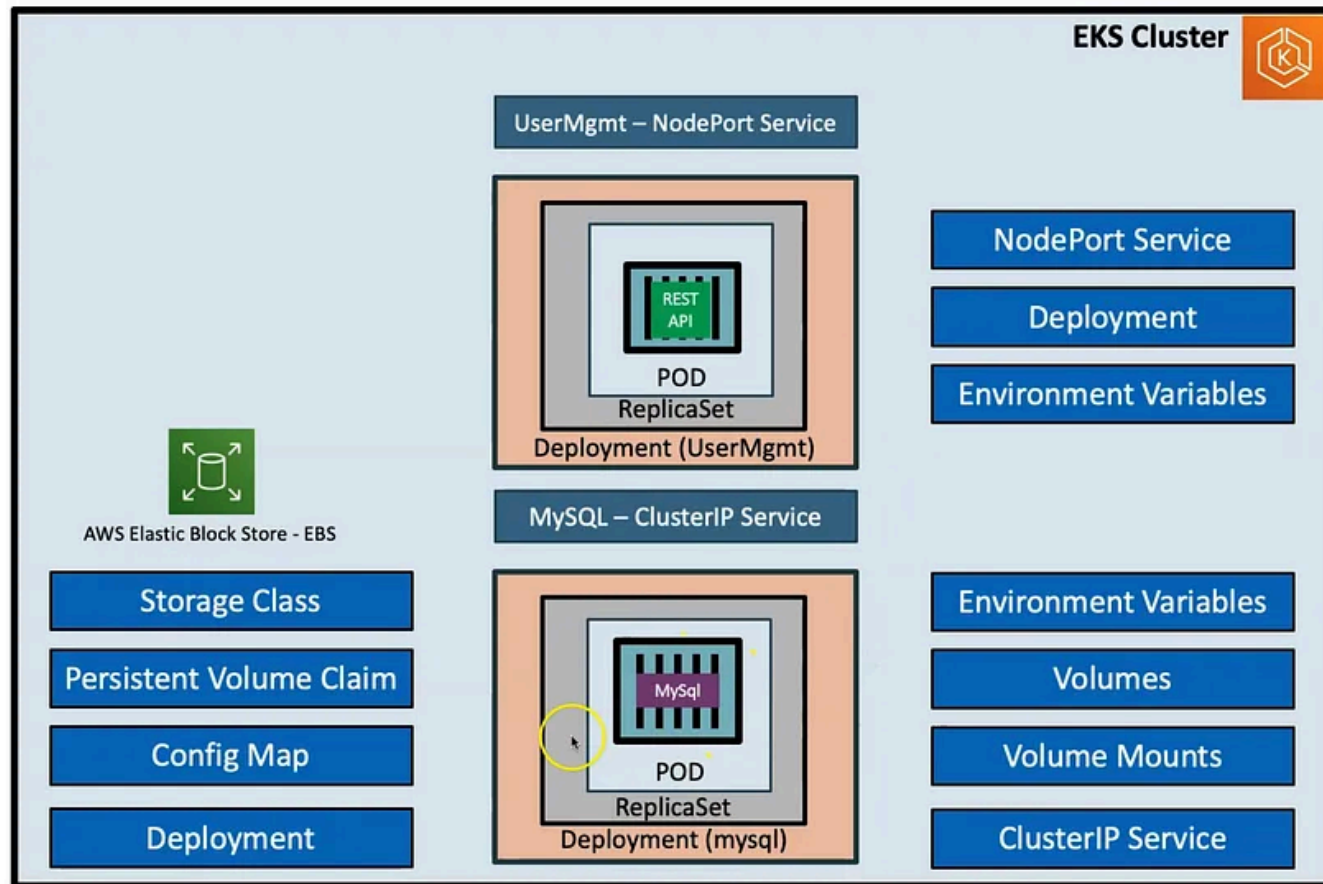


# kubernetes

## **Introduction**

The goal of this article is to learn Kubernetes by doing. Particularly, we're going to deploy the User Management REST API that utilizes MySQL DB.

Here's the architecture overview



## Namespace, LimitRange and ResourceQuota

- Namespaces allow to split-up resources into different groups.
- Resource names should be unique in a namespace.
- We can use namespaces to create multiple environments like dev, staging and production etc.

- Kubernetes will always list the resources from `default` namespace unless we provide exclusively from which namespace we need information from.
- Instead of specifying resources like `cpu` and `memory` in every container spec of a pod definition, we can provide the default CPU & Memory for all containers in a namespace using `LimitRange`.
- `ResourceQuota` allows us to specify the limits (in terms of CPU, Memory, and number of default kubernetes resources) on the namespace as a whole.

```
apiVersion: v1
kind: Namespace
metadata:
  name: dev
---
apiVersion: v1
kind: LimitRange
metadata:
  name: default-cpu-mem-limit-range
  namespace: dev
spec:
  limits:
    # every pod (container) in dev namespace won't consume more than this
    - default:
        memory: 512Mi
        cpu: 500m
```

```
# every pod (container) in dev namespace allocate at least this amount mem
defaultRequest:
  memory: 256Mi
  cpu: 250m
  type: Container
---
apiVersion: v1
kind: ResourceQuota
metadata:
  name: ns-resource-quota
  namespace: dev
spec:
  # namespace will take minimum of 1 vCPU and 1Gi of RAM
  # and won't consume more than 2 vCPUs and 2Gi of RAM.
  # We also specify the limits on number of particular k8s resources
  # that can be deployed in dev namespace
  hard:
    limits.cpu: 2
    limits.memory: 2Gi
    requests.cpu: 1
    requests.memory: 1Gi
    pods: 5
    configmaps: 5
    services: 5
    secrets: 5
    persistentvolumeclaims: 5
```

Verify the Resource Quotas and Limit Ranges

```
kubectl describe namespace dev
```

```
Name:      dev
Labels:    kubernetes.io/metadata.name=dev
Annotations: <none>
Status:    Active

Resource Quotas
Name:      ns-resource-quota
Resource   Used   Hard
-----
configmaps 2     5
limits.cpu  1     2
limits.memory 1Gi   2Gi
persistentvolumeclaims 1     5
pods        2     5
requests.cpu 500m   1
requests.memory 512Mi 1Gi
secrets     1     5
services    2     5

Resource Limits
Type      Resource  Min  Max  Default Request  Default Limit  Max Limit/Request Ratio
-----
Container cpu      -    -    250m             500m            -
Container memory -    -    256Mi            512Mi            -
```

## StorageClass

- Instead of creating **PersistentVolume**, that MySQL deployment will request with **PersistentVolumeClaim**, we'll allow Kubernetes to dynamically allocate persistent storage from attached EBS volumes to our EC2 instances that represent Worker Nodes of Kubernetes cluster.

- This is done via StorageClass resource that doesn't belong to any namespace.

```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: ebs-sc
provisioner: ebs.csi.aws.com # AWS EBS provision our persistent storage
volumeBindingMode: WaitForFirstConsumer # k8s won't allocate any volume unless t
```


## PersistentVolumeClaim

- StorageClass just allows us to utilize the underlying volume that is binded to our EC2 instance. In order for a pod to use storage, it needs to make a request for it specifying the desired amount


```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: ebs-mysql-pv-claim
  namespace: dev
spec:
  accessModes:
```

```
- ReadWriteOnce
storageClassName: ebs-sc # refer to the StorageClass we created previously
resources:
  requests:
    storage: 4Gi # specify the amount of volume we need
```

## ConfigMap

- Kubernetes resource for storing non-sensitive configuration data for our deployed microservices.
- 
- When we pass values to our environmental variables for pods, we refer to the keys mentioned in **ConfigMap**.

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: usermanagement-dbcreation-script
  namespace: dev
data:
  # this is gonna be a SQL-script that will configure our MySQL DB on the start.
  mysql_usermgmt.sql: |-
    DROP DATABASE IF EXISTS usermgmt;
    CREATE DATABASE usermgmt;
```





# MySQL Deployment

- Most of the time we don't deploy pods directly.
- We utilize the **Deployment** k8s resource for managing maintaining pods in a desired state.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: mysql
  namespace: dev
spec:
  # maintain just 1 replica of a pod
  replicas: 1
  strategy:
    # if we upgrade the pod, the old one will be first dropped
    # and afterwards a new one will spin up
    type: Recreate
  selector:
    # deployment will be responsible for the pods
    # with the labels app-mysql
    matchLabels:
      app: mysql
  template:
    metadata:
      # labels attached to the actual pod
      # should correspond to the one deployment specifies
      labels:
        app: mysql
    spec:
      containers:
```

```

- name: mysql
  image: mysql:5.6
  ports:
    - containerPort: 3306
      name: mysql
  # actual place within the file system of a pod
  # where we bind our volumes
  volumeMounts:
    - name: mysql-persistent-storage
      mountPath: /var/lib/mysql
    - name: usermgmt-dbcreation-script
      mountPath: /docker-entrypoint-initdb.d
  env:
    - name: MYSQL_ROOT_PASSWORD
      valueFrom:
        # value of password is a sensitive info
        # so it can't be stored in ConfigMap
        # we utilize another k8s resource - Secret
        secretKeyRef:
          name: mysql-secrets
          key: mysql-db-password
  # we list volumes that our pod will require
  volumes:
    # first is a request for the actual storage of a DB data
    - name: mysql-persistent-storage
      persistentVolumeClaim:
        claimName: ebs-mysql-pv-claim
    # second is the SQL-script we need to execute on a startup
    # that's stored in ConfigMap
    - name: usermgmt-dbcreation-script
      configMap:
        name: usermanagement-dbcreation-script

```

## Verify MySQL Deployment

```
# connect to our mysql pod
kubectl exec -it -n dev mysql-698ff4d9f8-dkmdr bash

# connect to the mysql server running within it
mysql -u root -pdbpassword11

# you should see our database we specified in configmap
mysql> show schemas;
```

## ClusterIP Service for MySQL Deployment

- Pods are naturally exposable and there're tons of reasons to restart them.
- Once the pod is restarted it changes its internal-IP. Hence other pods that were utilizing this IP directly will fail to communicate.
- We need something on top of a Pod for communication that won't change on Pod restarts.
- ClusterIP Service is exactly for that. We won't be able to access our Pod from outside of a k8s cluster, but internal communication between pods will happen through this service.

```
apiVersion: v1
kind: Service
metadata:
  name: mysql
  namespace: dev
spec:
  selector:
    app: mysql
  ports:
    - port: 3306
  clusterIP: None
```

## Verify the MySQL Deployment via service

```
# we specify host as the name of ClusterIP service
kubectl run -it --rm --image=mysql:5.6 --restart=Never \
  -n dev mysql-client01 -- mysql -h mysql -pdbpassword11
```

## Secret

- Just like a ConfigMaps, but for storing a sensitive information within a Kubernetes cluster.

```
apiVersion: v1
kind: Secret
metadata:
  name: mysql-secrets
  namespace: dev
type: Opaque
data:
  # we specify key-value pairs of sensitive info
  # where values are not stored in a raw format, but base64-encoded
  mysql-db-user: cm9vdA== # base64 encoded root
  mysql-db-password: ZGJwYXNzd29yZDEx
  mysql-db-hostname: bXlzcWw= # base64 encoded mysql value
  mysql-db-host: bXlzcWw=
  mysql-db-name: dXNlcm1nbXQ= # base64 encoded usermgmt value
```

## User Management Deployment

- This is our REST API that will communicate with MySQL Deployment for persistently storing information about users.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: usermgmt
  namespace: dev
spec:
  replicas: 1
```

```

strategy:
  type: Recreate
selector:
  matchLabels:
    app: usermgmt
template:
  metadata:
    labels:
      app: usermgmt
  spec:
    # specify container that spin up first before our main app process
    # it'll run in a loop checking the availability of a MySQL deployment
    # once MySQL respond back, we stop the init container and start the main o
    # this used to prevent unnecessary pod restarts, when depended pod isn't u
    initContainers:
      - name: init-db
        image: busybox:1.31
        command: ['sh', '-c', 'echo -e "Checking for the availability of MySQL
containers:
      - name: usermgmt
        image: stacksimplify/kube-usermanagement-microservice:1.0.0
        # bunch of env vars that aren't directly specified here
        # but actually stored in a secret and referenced here
        env:
          - name: DB_HOSTNAME
            valueFrom:
              secretKeyRef:
                name: mysql-secrets
                key: mysql-db-hostname

          - name: DB_PORT
            value: "3306"

          - name: DB_NAME
            valueFrom:
              secretKeyRef:

```

```
    name: mysql-secrets
    key: mysql-db-name

- name: DB_USERNAME
  valueFrom:
    secretKeyRef:
      name: mysql-secrets
      key: mysql-db-user

- name: DB_PASSWORD
  valueFrom:
    secretKeyRef:
      name: mysql-secrets
      key: mysql-db-password

ports:
  - containerPort: 8095
# once 60 seconds after the start of a container passes we run
# a check every 10 seconds to make sure the process is up and running
# without being stuck in any kind of a deadlock
livenessProbe:
  exec:
    command:
      - /bin/sh
      - -c
      - nc -z localhost 8095
  initialDelaySeconds: 60
  periodSeconds: 10
# this check that process is not only up and running
# but also accepts the HTTP traffic
readinessProbe:
  httpGet:
    path: /usermgmt/health-status
    port: 8095
  initialDelaySeconds: 60
  periodSeconds: 10
```

## NodePort Service for User Management

- We want to be able to access our User Management Deployment from outside of a k8s cluster.
- ClusterIP Service isn't suitable for it.
- NodePort is a simplest way to expose our Deployment to outside world. It basically opens a specific port on our Worker Nodes and all traffic that comes to it will be redirected to this service.
- Don't forget to add new inbound rules for Security Groups attached to EC2 instances that represent our Node Group. They need to allow any traffic to the port you specify in a **nodePort** key.

```
apiVersion: v1
kind: Service
metadata:
  name: mysql-svc
  namespace: dev
  labels:
    app: usermgmt
spec:
```



```

type: NodePort
ports:
  - port: 8095
    targetPort: 8095
    nodePort: 31231 # we can drop it and k8s will allocate port randomly in a
# specify to which pods traffic will be forwarded based on the labels of a pod
selector:
  app: usermgmt

```

## Verify the whole deployment

Run the `kubectl get all -n dev` to list all the deployed resources in our dev namespace

```

konstantinmogilevskii@Konstantins-MacBook-Pro mysql % kubectl get all -n dev
NAME                                READY    STATUS    RESTARTS   AGE
pod/mysql-698ff4d9f8-dkmdr          1/1      Running   0           62m
pod/mysql-client                    0/1      Completed 0           42m
pod/usermgmt-856c7c65b4-l4ccb       1/1      Running   0           7m45s

NAME                                TYPE          CLUSTER-IP    EXTERNAL-IP    PORT(S)          AGE
service/mysql                        ClusterIP     None           <none>          3306/TCP         58m
service/mysql-svc                    NodePort      10.100.144.33 <none>          8095:32396/TCP   5m

NAME                                READY    UP-TO-DATE   AVAILABLE   AGE
deployment.apps/mysql               1/1      1             1           62m
deployment.apps/usermgmt            1/1      1             1           7m46s

NAME                                DESIRED    CURRENT   READY   AGE
replicaset.apps/mysql-698ff4d9f8    1          1         1       62m
replicaset.apps/usermgmt-856c7c65b4 1          1         1       7m46s

```

As you can see **mysql-svc** Service forwards traffic from the port 32396 to 8095 of a Pod that runs User Management system. That's because in the actual steps I commented the nodePort and allowed k8s to pick one randomly for me.

Let's check how many EC2 instance resources are consumed by our deployments:

```
kubectl get nodes  
  
kubectl describe node <node-name>
```

Namespace	Name	CPU Requests	CPU Limits	Memory Requests	Memory Limits	Age
dev	usermgmt-856c7c65b4-l4ccb	250m (12%)	500m (25%)	256Mi (7%)	512Mi (15%)	21m
kube-system	aws-node-jbh8p	50m (2%)	0 (0%)	0 (0%)	0 (0%)	142m
kube-system	coredns-54d6f577c6-l59h9	100m (5%)	0 (0%)	70Mi (2%)	170Mi (5%)	179m
kube-system	coredns-54d6f577c6-w2wkv	100m (5%)	0 (0%)	70Mi (2%)	170Mi (5%)	179m
kube-system	ebs-csi-controller-9f85688cd-snf2d	60m (3%)	0 (0%)	240Mi (7%)	1536Mi (46%)	134m
kube-system	ebs-csi-node-dx59m	30m (1%)	0 (0%)	120Mi (3%)	768Mi (23%)	134m
kube-system	kube-proxy-2l9d9	100m (5%)	0 (0%)	0 (0%)	0 (0%)	142m

One node will have MySQL deployment and other the UserManagement Deployment. No matter what the process is doing it still consumes at minimum the 250m vCPU and 256Mi of RAM.

Once you altered the Security Group Inbound Rules for our Node Group, we can check the API with several calls:

GET request to `http://34.229.193.196:32396/usermgmt/health-status`

POST request to `http://34.229.193.196:32396/usermgmt/user` to create a user and s  
Body JSON can be

```
{  
  "username": "admin1",  
  "email": "dkalyanreddy@gmail.com",  
  "role": "ROLE_ADMIN",  
  "enabled": true,  
  "firstname": "fname1",  
  "lastname": "lname1",  
  "password": "Pass@123"  
}
```

GET to `http://34.229.193.196:32396/usermgmt/users` to list all the present `users`

Thanks!

Kubernetes

Microservices

Aws Eks



## Written by Konstantin Mogilevskii


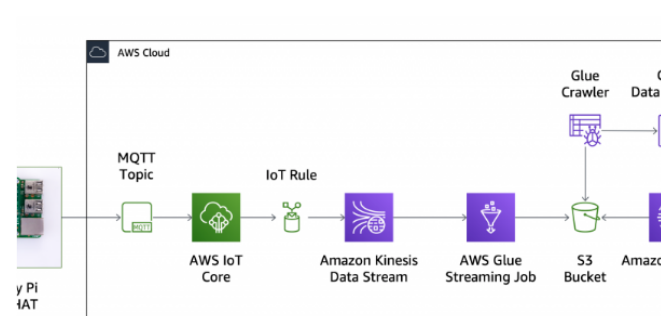
230 Followers · Writer for Dev Genius

Senior Data Engineer [www.linkedin.com/in/kmogilevskii](https://www.linkedin.com/in/kmogilevskii)

Follow



### More from Konstantin Mogilevskii and Dev Genius



 Konstantin Mogilevskii in Dev Genius

## Serverless Streaming ETL with AWS Glue

Introduction

★ · 7 min read · May 9, 2024



 Ani Talakhadze in Dev Genius

## Top 10 Essential Stream Operations in Java

Boost Your Java Efficiency With Examples, Exercises And Fun Facts

9 min read · Apr 21, 2024




 Tomas Svojanovsky in Dev Genius

## Pytest Fixtures: Your Secret Weapon for Writing Powerful Tests

Exploring Pytest Fixtures: Setup, Teardown, Scopes, and Best Practices

★ · 6 min read · Apr 15, 2024



 Konstantin Mogilevskii in Dev Genius

## Practical Data Modeling with Data Vault 2.0 and Vertica

Introduction

★ · 6 min read · Mar 23, 2024

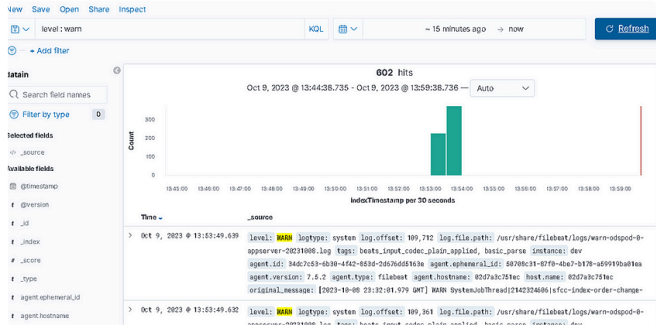


[See all from Konstantin Mogilevskii](#)

[See all from Dev Genius](#)

---

## Recommended from Medium



Yassine Elmahi

## Using a Dockerized ELK-F stack to consume and analyse logs

What is the ELK stack?

8 min read · 3 days ago



91

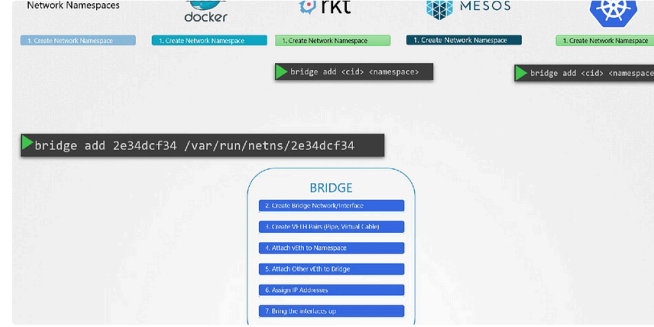


## Lists



## Natural Language Processing

1456 stories · 966 saves



Devendra Johari in DevOps.dev

## Networking in Kubernetes

Kubernetes

18 min read · May 1, 2024



 Tanat Lokejaroenlarb in NonTechCompany

## ArgoCd ApplicationSet is more practical in version v2.9

3 min read · Apr 27, 2024




25



 Tony



 Tamer Benhassan



## 7 Best Practices for Implementing Security Contexts in Kubernetes

Kubernetes security contexts are a fundamental part of safeguarding Kuberne...

3 min read · May 9, 2024



53



 Spacelift in Spacelift

## Most Popular Jenkins Alternatives



# Streamlit—Next-Gen Web App Development Framework

Usually, before diving into learning a new framework, people typically research what...

🌟 · 7 min read · 4 days ago



26



1



Jenkins was a pioneer in the CI/CD space, but today, there are many alternatives available....

9 min read · Apr 22, 2024



9



1



See more recommendations